

▼ PRACTICE-4:

Design a model to predict the housing price from Boston Dataset using Linear Regression

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import pandas as pd
4 import seaborn as sns
5 %matplotlib inline
```

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Mounted at /content/drive

```
1 import io
2 import pandas as pd
3 df=pd.read_csv('/content/drive/MyDrive/Data Sets/BostonHousing.csv')
4 df.head()
```

```

      crim    zn  indus  chas   nox    rm  age   dis  rad  tax  ptratio    b  lstat  medv
0  0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900  1  296    15.3  396.90   4.98  24.0
1  0.02731   0.0   7.07    0  0.469  6.421  78.9  4.9671  2  242    17.8  396.90   9.14  21.6
2  0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671  2  242    17.8  392.83   4.03  34.7
3  0.03237   0.0   2.18    0  0.458  6.998  45.8  6.0622  3  222    18.7  394.63   2.94  33.4
4  0.06905   0.0   2.18    0  0.458  7.147  54.2  6.0622  3  222    18.7  396.90   5.33  36.2
```

```
1 df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  ---
 0   crim      506 non-null    float64
 1   zn        506 non-null    float64
 2   indus     506 non-null    float64
 3   chas      506 non-null    int64
 4   nox       506 non-null    float64
 5   rm        506 non-null    float64
 6   age       506 non-null    float64
 7   dis       506 non-null    float64
 8   rad       506 non-null    int64
 9   tax       506 non-null    int64
10  ptratio   506 non-null    float64
11  b         506 non-null    float64
12  lstat     506 non-null    float64
13  medv     506 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB
```

```
1 df.describe()
```

```


      crim    zn  indus  chas   nox    rm  age   dis  rad  tax  ptratio    b  lstat  medv
count  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000
mean    3.613524  11.363636  11.136779   0.069170   0.554695   6.284634  68.574901   3.795043   9.549407  408.237154  18.455534  356.674032  12.653063  22.532806
std     8.601545  23.322453   6.860353   0.253994   0.115878   0.702617  28.148861   2.105710   8.707259  168.537116   2.164946   91.294864   7.141062   9.197104
min     0.006320   0.000000   0.460000   0.000000   0.385000   3.561000   2.900000   1.129600   1.000000  187.000000  12.600000   0.320000   1.730000   5.000000
25%     0.082045   0.000000   5.190000   0.000000   0.449000   5.885500  45.025000   2.100175   4.000000  279.000000  17.400000  375.377500   6.950000  17.025000
50%     0.256510   0.000000   9.690000   0.000000   0.538000   6.208500  77.500000   3.207450   5.000000  330.000000  19.050000  391.440000  11.360000  21.200000
75%     3.677083  12.500000  18.100000   0.000000   0.624000   6.623500  94.075000   5.188425  24.000000  666.000000  20.200000  396.225000  16.955000  25.000000
```

```

1 #let's plot all the columns to look at their distributions
2 for i in df.columns:
3     plt.figure(figsize=(7,4))
4     sns.histplot(data=df,x=i,kde=True)
5     plt.show()
```

```

1 plt.figure(figsize=(12,8))
2 cmap=sns.diverging_palette(230,20,as_cmap=True)
3 sns.heatmap(df.corr(),annot=True,fmt='.2f',cmap=cmap)
4 plt.show()
```



```

1 Y=df['medv']
2 X=df.drop(columns={'medv'})

1 #splitting the data in 70:30 ratio of train to test data
2 from sklearn.model_selection import train_test_split
3 X_train,X_test,y_train,y_test=train_test_split(X,Y,test_size=0.30,random_state=1)

```

```

1 print(X_train.shape)
2 print(X_test.shape)
3 print(y_train.shape)
4 print(y_test.shape)

```

```

(354, 13)
(152, 13)
(354,)
(152,)

```

```

1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import mean_squared_error
3 lin_model=LinearRegression()
4 lin_model.fit(X_train,y_train)

```

```
LinearRegression
```

```

1 b=lin_model.coef_
2 m=lin_model.intercept_

```

```
1 b
```

```

array([-9.85424717e-02,  6.07841138e-02,  5.91715401e-02,  2.43955988e+00,
        -2.14699650e+01,  2.79581385e+00,  3.57459778e-03, -1.51627218e+00,
         3.07541745e-01, -1.12800166e-02, -1.00546640e+00,  6.45018446e-03,
        -5.68834539e-01])

```

```
1 m
```

```
46.396493871823864
```

```

1 from sklearn.metrics import r2_score
2 #model evaluation for training set
3 y_train_predict=lin_model.predict(X_train)
4 rmse=(np.sqrt(mean_squared_error(y_train,y_train_predict)))
5 r2=r2_score(y_train,y_train_predict)
6
7 print("The model performance for training set")
8 print("-----")
9 print('RMSE is {}'.format(rmse))
10 print('R2 score is {}'.format(r2))
11 print("\n")

```

```

The model performance for training set
-----
RMSE is 4.849055005805464
R2 score is 0.7103879080674731

```

```

1 from sklearn.metrics import r2_score
2 #model evaluation for training set '
3 y_test_predict=lin_model.predict(X_test)
4 rmse=(np.sqrt(mean_squared_error(y_test,y_test_predict)))
5 r2=r2_score(y_test,y_test_predict)
6
7 print("The model performance for training set")
8 print("-----")
9 print('RMSE is {}'.format(rmse))
10 print('R2 score is {}'.format(r2))
11 print("\n")
12
13

```

```

The model performance for training set
-----
RMSE is 4.453237437198149
R2 score is 0.7836295385076292

```

HW:

1. Write the regression equation $y_{pred} = m + b_0x_1 + b_1x_2 + \dots + b_{12}x_{13}$
2. Select one datasample from the dataset
3. find y_{pred} value
4. find the absolute error for y_{pred} and y_{actual}

```

1 l=[]
2 l=list(df.iloc[3])
3 b=list(b)
4 print(l)
5 print(b)
6
7

```

```

[0.03237, 0.0, 2.18, 0.0, 0.458, 6.998, 45.8, 6.0622, 3.0, 222.0, 18.7, 394.63, 2.94, 33.4]
[-0.09854247168320542, 0.060784113842432486, 0.05917154006562941, 2.4395598751549907, -21.469964963213876, 2.7958138545358477, 0.0035745977759409353, -1.516272184239845, 0.307541745464

```

```

1 ypred=m
2 for i in range(len(b)):
3     ypred=ypred+(b[i]*l[i])
4 print(ypred)
5

```